

CAB401 High Performance and Parallel Computing

**Project Assignment
(DigitalMusicAnalysis)**

Queensland University of Technology

Pattarachai Roongsritong n10548467

Table of Contents

1.	<i>Introduction</i>	2
1.1	Digital Music Analysis program overview	2
2.	<i>Accessed Tools</i>	3
2.1	Hardware	3
2.2	Software	3
3.	<i>Parallelising Implementation</i>	3
3.1	Preliminary profiling	3
3.2	timefreq class	4
3.3	onsetDetection method	8
4.	<i>Results</i>	10
4.1	Sequential version	10
4.2	Parallelised version	10
4.3	Outputs	12
5.	<i>Reflection</i>	14
6.	<i>Appendices</i>	14
6.1	Appendix A	14
6.2	Appendix B	15

1. Introduction

This report is going to demonstrate an attempt to enhance the performance of a program called “digitalMusicAnalysis” which is a program to help rookie violin players get their performance feedback in terms of correctness. This project aims to achieve through leveraging parallel computing tools and techniques.

1.1 Digital Music Analysis program overview

As mentioned earlier, the program is used to give a violin player feedback of his/her recorded performance whether the high (sharp) or the low (flat) notes are played correctly. The program takes a WAV file input in companion with an XML file input as a reference then visualises sound frequencies and octaves highlighting the difference from correct octaves according to the reference. The UML class diagram is shown in the **figure 1**.

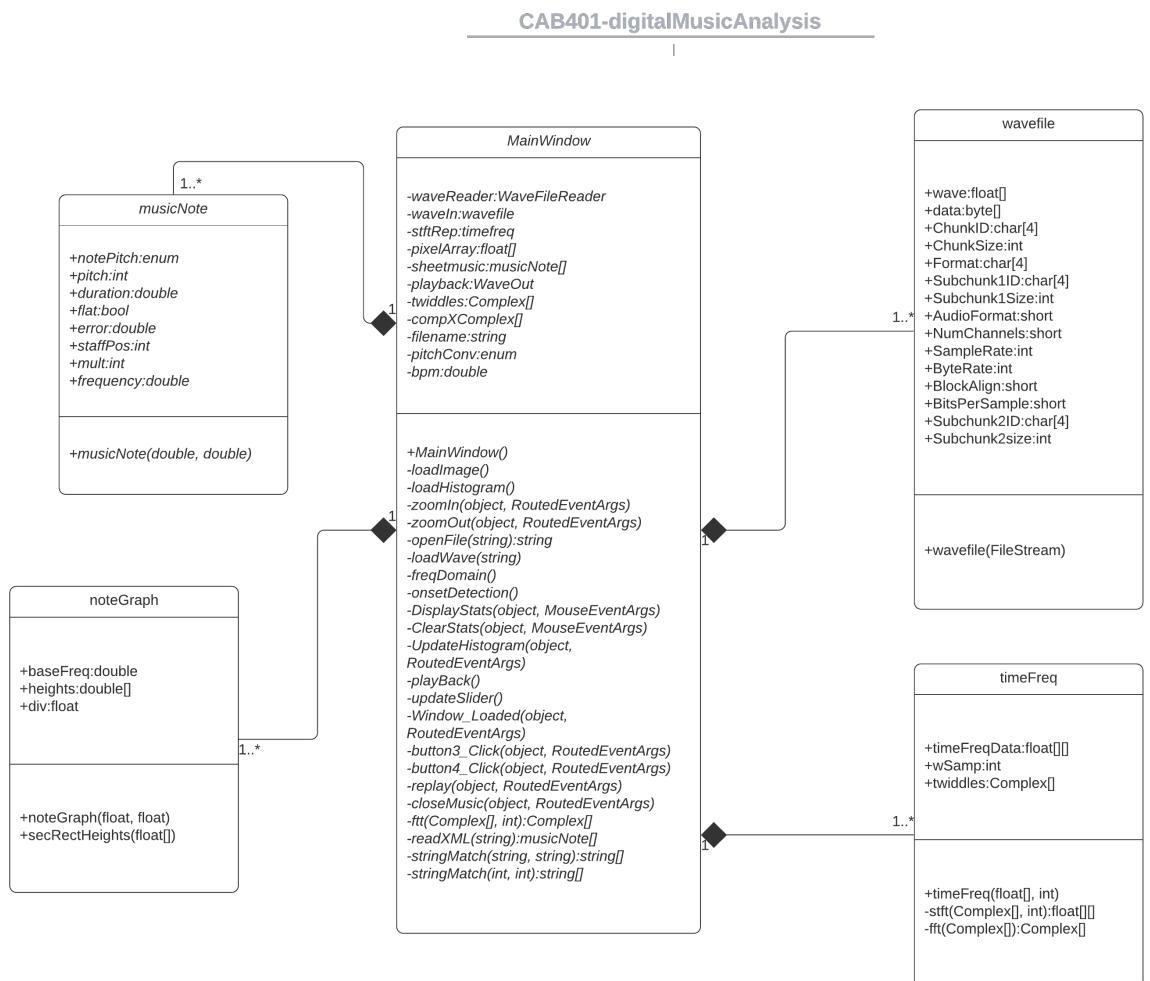


Fig.1 UML class diagram of the digitalMusicAnalysis program

2. Accessed Tools

2.1 Hardware

Since I can only access to a MAC, The QUT virtual machine is used via the VMware Horizontal Client to parallelise the application which has the following specifications.

Processor - Intel Xeon Processor E5-2687W v4

Clock speed – 3.00 GHz

Cache – 30MB

The number of cores – 12

The number of threads – 24

Installed memory – 16.0 GB

2.2 Software

The original program was developed in C# thus the Microsoft Visual Studio 2019 is used to achieve parallelism by the Task Parallel Library (TPL) as well as the built-in performance profiler to analyse the CPU usage data. Moreover, the built-in C# library is used to track the running time.

3. Parallelising Implementation

3.1 Preliminary profiling

The performance profiler is used to determine how much CPU resources are spent in each program component. According to the result as shown in the **figure 2**, it suggests that the fff functions spend the most CPU resources. The functions appear in two spots of the program which are the MainWindow class and the timefreq class. Moreover, the hot path reveals that the two fff functions are invoked through the freqDomain (line 48) and the onsetDetection (line 51) functions. Subsequently, when taking a closer look at the freqDomain function, the function spends most of its time instantiating an object from timeFreq class (DigitalMusicAnalysis.timefreq..ctor(float[]32, int)). As a result, our main investigating spots are the timeFreq class and the onsetDetection function.

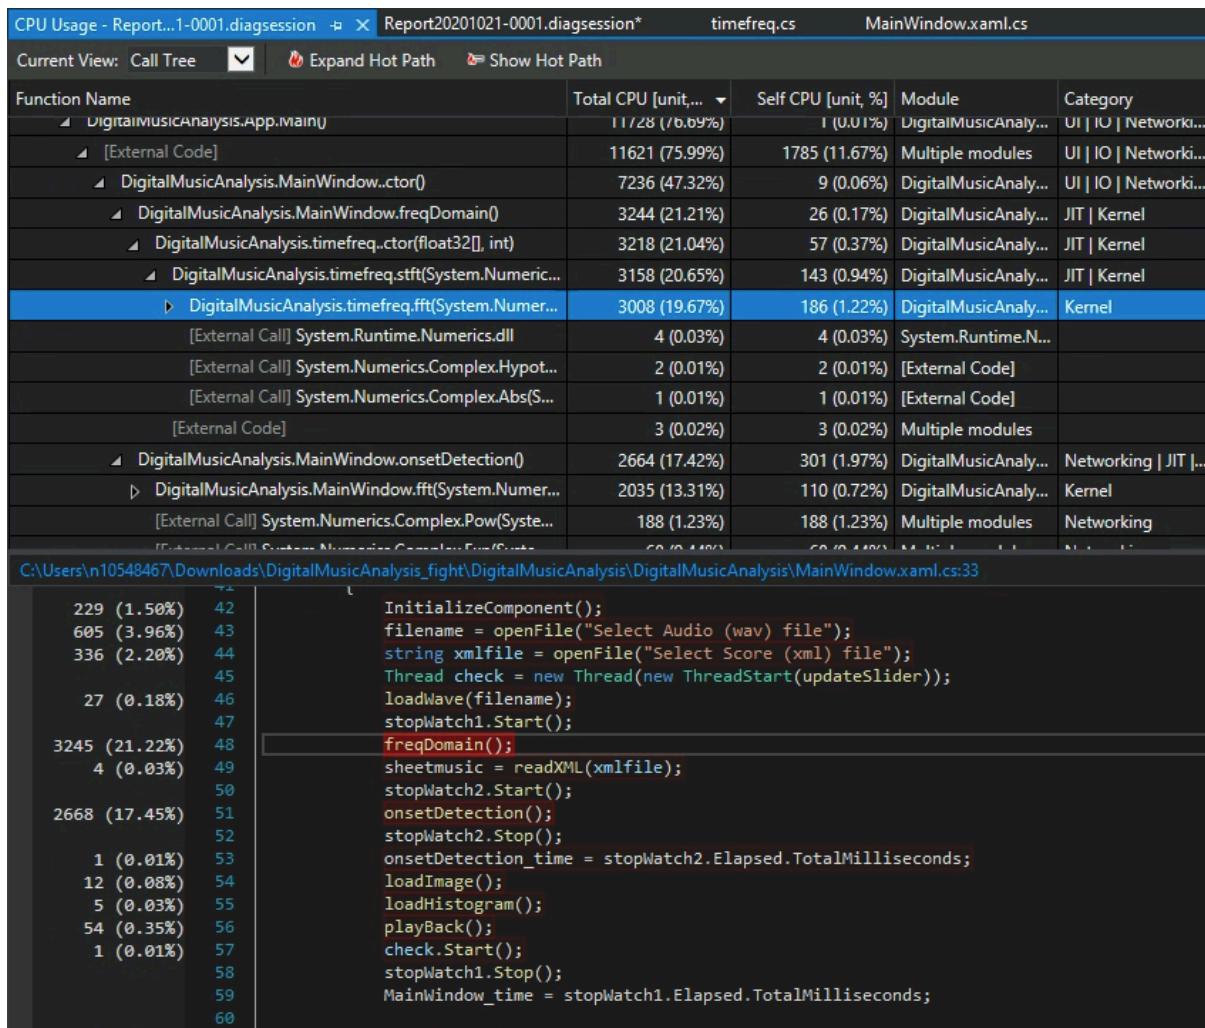


Fig.2 Profiling result of the sequential version program

3.2 timefreq class

Initial analysis

According to the profiler shown in **figure 3, 4, and 5**, within the constructor of this class (line 285 of MainWindow.xaml.cs), the main work is done by calling the stff function (line 53 of timefreq.cs) to calculate a two-dimensional array of floating-point numbers which requires to call the fff function (line 87 of timefreq.cs) a number of times. In order to safely parallelise, there are two data dependencies to be considered which are the two temporary arrays of complex numbers (line 76 and 77 of timefreq.cs). Additionally, the profiler also suggests that the loop involving with a calculation of twiddle numbers (**figure 6**) and the loop calculating the Y array (**figure 7**) have some sample counts, so they are interesting to be parallelised as well because they carry no dependency.

```

C:\Users\n10548467\Downloads\DigitalMusicAnalysis_fight\DigitalMusicAnalysis\DigitalMusicAnalysis\MainWindow.xaml.cs:284
    }
    // Transforms data into Time-Frequency representation
283     private void freqDomain()
284     {
285         stftRep = new timefreq(waveIn.wave, 2048);
286         pixelArray = new float[stftRep.timeFreqData[0].Length * stftRep.wSamp / 2];
287         for (int jj = 0; jj < stftRep.wSamp / 2; jj++)
288         {
289             for (int ii = 0; ii < stftRep.timeFreqData[0].Length; ii++)
290             {
291                 pixelArray[jj * stftRep.timeFreqData[0].Length + ii] = stftRep.timeFreqData[jj][ii];
292             }
293         }
294     }
295 }
296

```

Fig.3 Using timefreq constructor in freqDomain()

```

C:\Users\n10548467\Downloads\DigitalMusicAnalysis_fight\DigitalMusicAnalysis\DigitalMusicAnalysis\timefreq.cs:12
    }
    int cols = 2 * nearest / wSamp;
    for (int jj = 0; jj < wSamp / 2; jj++)
    {
        timeFreqData[jj] = new float[cols];
    }
    stopWatch.Start();
    timeFreqData = stft(compX, wSamp);
    stopWatch.Stop();
    stft_time = stopWatch.Elapsed.TotalMilliseconds;
}
float[][] stft(Complex[] x, int wSamp)
{
    int ii = 0;
    int jj = 0;
    int kk = 0;
    int ll = 0;
}

```

Fig.4 Invoking of stft function

```

58     float[][] stft(Complex[] x, int wSamp)
59     {
60         int ii = 0;
61         int jj = 0;
62         int kk = 0;
63         int ll = 0;
64         int N = x.Length;
65         float fftMax = 0;
66
67         float[][] Y = new float[wSamp / 2][];
68
69
70         for (ll = 0; ll < wSamp / 2; ll++)
71         {
72             Y[ll] = new float[2 * (int)Math.Floor((double)N / (double)wSamp)];
73         }
74
75         Complex[] temp = new Complex[wSamp];
76         Complex[] tempFFT = new Complex[wSamp];
77
78         for (ii = 0; ii < 2 * Math.Floor((double)N / (double)wSamp) - 1; ii++)
79         {
80
81             for (jj = 0; jj < wSamp; jj++)
82             {
83                 temp[jj] = x[ii * (wSamp / 2) + jj];
84             }
85
86             tempFFT = fft(temp);
87
88             for (kk = 0; kk < wSamp / 2; kk++)
89             {
90                 Y[kk][ii] = (float)Complex.Abs(tempFFT[kk]);
91             }
92
93             if (Y[kk][ii] > fftMax)
94             {
95                 fftMax = Y[kk][ii];
96             }
97         }
98

```

Fig.5 Invoking fft function from stft function

```

20      twiddles = new Complex[wSamp];
21  □    for (ii = 0; ii < wSamp; ii++)
22  {
23      double a = 2 * pi * ii / (double)wSamp;
24      twiddles[ii] = Complex.Pow(Complex.Exp(-i), (float)a);
25  }
26

```

Fig.6 A loop of Twiddle numbers calculation

```

70
71  □    for (ll = 0; ll < wSamp / 2; ll++)
72  {
73      Y[ll] = new float[2 * (int)Math.Floor((double)N / (double)wSamp)];
74  }
75

```

Fig.7 A loop of Y array calculation

Approach

Due to its dependencies, the whole code-block of ii loop containing the fff function is performed on its own thread to achieve parallelism by creating a new method as the **figure 8** called stftOnThread to run on an individual thread as the **figure 9**. The loops for calculating twiddle values and Y array are also implemented parallelisation by the Parallel.For method which the iterations are able to run in parallel and the maximum number of parallel tasks can be set according to our needs as shown in **figure 10 and 11**.

```

C:\Users\n10548467\Downloads\DigitalMusicAnalysis_fight\DigitalMusicAnalysis - Parallelised\DigitalMusicAnalysis\timefreq.cs:25
    118     public void stftOnThread(object threadID)
    119     {
    120         int id = (int)threadID;
    121         int start = id * chunk_size;
    122
    123         Complex[] temp = new Complex[wSamp];
    124         Complex[] tempFFT = new Complex[wSamp];
    125
    126         for (int ii = start; ii < Math.Min(start + chunk_size, array_length - 1); ii++)
    127         {
    128             for (int jj = 0; jj < wSamp; jj++)
    129             {
    130                 temp[jj] = newX[ii] * (wSamp / 2) + jj;
    131             }
    132
    133             tempFFT = fft(temp);
    134
    135             for (int kk = 0; kk < wSamp / 2; kk++)
    136             {
    137                 Y[kk][ii] = (float)Complex.Abs(tempFFT[kk]);
    138
    139                 if (Y[kk][ii] > fftMax)
    140                 {
    141                     fftMax = Y[kk][ii];
    142                 }
    143             }
    144         }
    145     }

```

Fig.8 stftOnThread Method

```

    73     float[][] stft(Complex[] x, int wSamp)
    74     {
    75         N = x.Length;
    76         fftMax = 0;
    77         newX = x;
    78         array_length = 2 * (int)Math.Floor(N / (double)wSamp)-1;
    79         chunk_size = (array_length + MainWindow.NUM_THREADS - 1) / MainWindow.NUM_THREADS;
    80
    81         Y = new float[wSamp / 2][];
    82
    83
    84
    85         Parallel.For(0, wSamp / 2, new ParallelOptions { MaxDegreeOfParallelism = MainWindow.NUM_THREADS }, ll =>
    86         {
    87             Y[ll] = new float[2 * (int)Math.Floor((double)N / (double)wSamp)];
    88         });
    89
    90
    91         Thread[] fft_thread = new Thread[MainWindow.NUM_THREADS];
    92
    93         for (int i = 0; i < MainWindow.NUM_THREADS; i++)
    94         {
    95             fft_thread[i] = new Thread(stftOnThread);
    96             fft_thread[i].Start(i);
    97         }
    98         for (int j = 0; j < MainWindow.NUM_THREADS; j++)
    99         {
    100             fft_thread[j].Join();
    101         }
    102
    103
    104
    105         for (int ii = 0; ii < 2 * Math.Floor((double)N / (double)wSamp) - 1; ii++)
    106         {
    107             for (int kk = 0; kk < wSamp / 2; kk++)
    108             {
    109                 Y[kk][ii] /= fftMax;
    110             }
    111         }
    112     }

```

Fig.9 Parallelising stft function by implementing stftOnThread

```

    34
    35     twiddles = new Complex[wSamp];
    36     Parallel.For(0, wSamp, MainWindow.parallel_options, ii =>
    37     {
    38         double a = 2 * pi * ii / (double)wSamp;
    39         twiddles[ii] = Complex.Pow(Complex.Exp(-i), (float)a);
    40     });

```

Fig.10 Parallelising a loop of Twiddle numbers calculation

```

    84
    85     Parallel.For(0, wSamp / 2, new ParallelOptions { MaxDegreeOfParallelism = MainWindow.NUM_THREADS }, ll =>
    86     {
    87         Y[ll] = new float[2 * (int)Math.Floor((double)N / (double)wSamp)];
    88     });

```

Fig.11 Parallelising a loop of Y array calculation

3.3 onsetDetection method

Initial analysis

As discussed in the preliminary profiling section, this method involves in calling the fff function multiple times and when we take a look at the code block, it also has its own expensive calculations which can be listed as follows:

- HFC array calculation loop

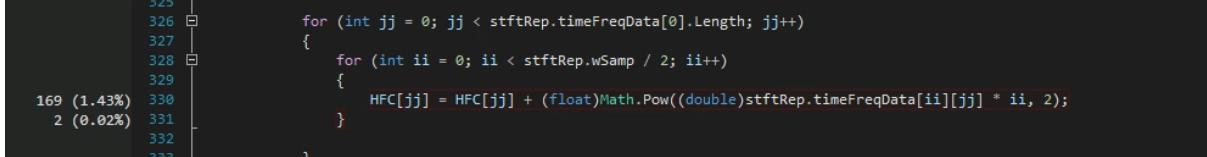


Fig.12 HFC array calculation loop (sequential)

- twiddles array calculation loop
- complex number array calculation loop which is an argument for the fff function

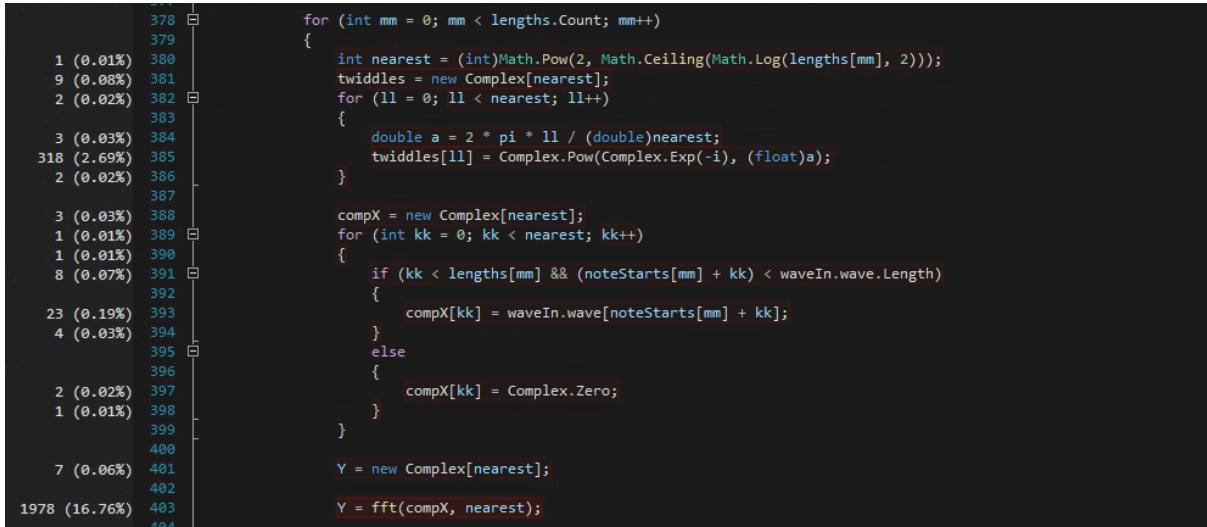


Fig.13 Twiddles array and complex number array calculation loops

- absolute values of Y array calculation loop

In order to enhance the performance of this method, the aforementioned loops will be given parallelism attempts.

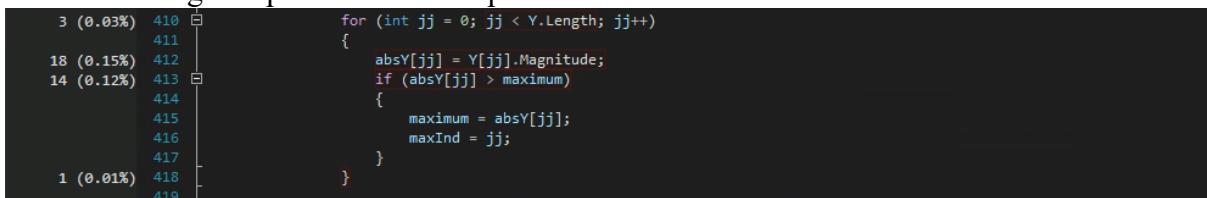


Fig.14 Absolute values of Y array calculation loop

Approach

- HFC array calculation loop: It seems that the structure of the loop allows temporal locality because the outer loop index jj is used to read and write the HFC array and read the timeFreqData array same elements multiple times within the inner loop ii which is likely to result in cache hit without a loop transformation; however, the HFC array calculation loop carries no

dependency so it is safe to parallelise which can potentially save some computing resources.

```

24 (0.14%) 334 Parallel.For(0, stftRep.timeFreqData[0].Length, parallel_options, jj =>
335 {
336     for (int ii = 0; ii < stftRep.wSamp / 2; ii++)
337     {
338         HFC[jj] = HFC[jj] + (float)Math.Pow((double)stftRep.timeFreqData[ii][jj] * ii, 2);
339     }
340 });
341

```

Fig.15 Parallelised HFC loop

- twiddles array and complex number array calculation loops: These two calculations can be parallelised by using a loop fusion since both loops keep running until the index is less than the variable called nearest and still can preserve the same data dependency (the variable a) and control dependency (if-else statement). Unfortunately, the attempt to parallelise the fff function calling is fail despite having done different ways of restructuring at the mm loop because its arguments (compX and nearest) carry data dependencies.

```

383
384
385 for (int mm = 0; mm < lengths.Count; mm++)
386 {
387     int nearest = (int)Math.Pow(2, Math.Ceiling(Math.Log(lengths[mm], 2)));
388     twiddles = new Complex[nearest];
389     compX = new Complex[nearest];
390
391     Parallel.For(0, nearest, parallel_options, zz =>
392     {
393         double a = 2 * pi * zz / (double)nearest;
394         twiddles[zz] = Complex.Pow(Complex.Exp(-i), (float)a);
395
396         if (zz < lengths[mm] && (noteStarts[mm] + zz) < waveIn.wave.Length)
397         {
398             compX[zz] = waveIn.wave[noteStarts[mm] + zz];
399         }
400         else
401         {
402             compX[zz] = Complex.Zero;
403         }
404     });
405
406     Y = new Complex[nearest];
407
408     Y = fft(compX, nearest);
409

```

Fig.16 Parallelised Twiddles array and complex number array calculation loops

- absolute values of Y array calculation loop: The operation which has a certain sample counts in this loop is assigning absolute values of Y array, so the loop is transformed by a loop distribution and parallelised.

```

48 (0.28%) 415 Parallel.For(0, Y.Length, parallel_options, jj =>
416 {
417     absY[jj] = Y[jj].Magnitude;
418 });
419 for (int jj = 0; jj < Y.Length; jj++)
420 {
421     if (absY[jj] > maximum)
422     {
423         maximum = absY[jj];
424         maxInd = jj;
425     }
426 }
427

```

Fig.17 Parallelised absolute values of Y array calculation loop

4. Results

The Stopwatch class is used to measure the time used to execute three functions including MainWindow representing total runtime of the program, stft, and onsetDetection program which the latter two are our hot spot to attempt parallelisation.

4.1 Sequential version

The time taken of each function can be summarised in the **table 1**.

Method name	time (ms)
MainWindow	3642.3662
stft	1754.5105
onsetDetection	1618.3864

Table 1. Spent time on sequential version

4.2 Parallelised version

In the **figure 18**, it can be seen that a profiling result has changed from the sequential version as the freqDomain function does not spend as much CPU resources as in prior profiling result. However, the onsetDetection function still consumes significant CPU resources.

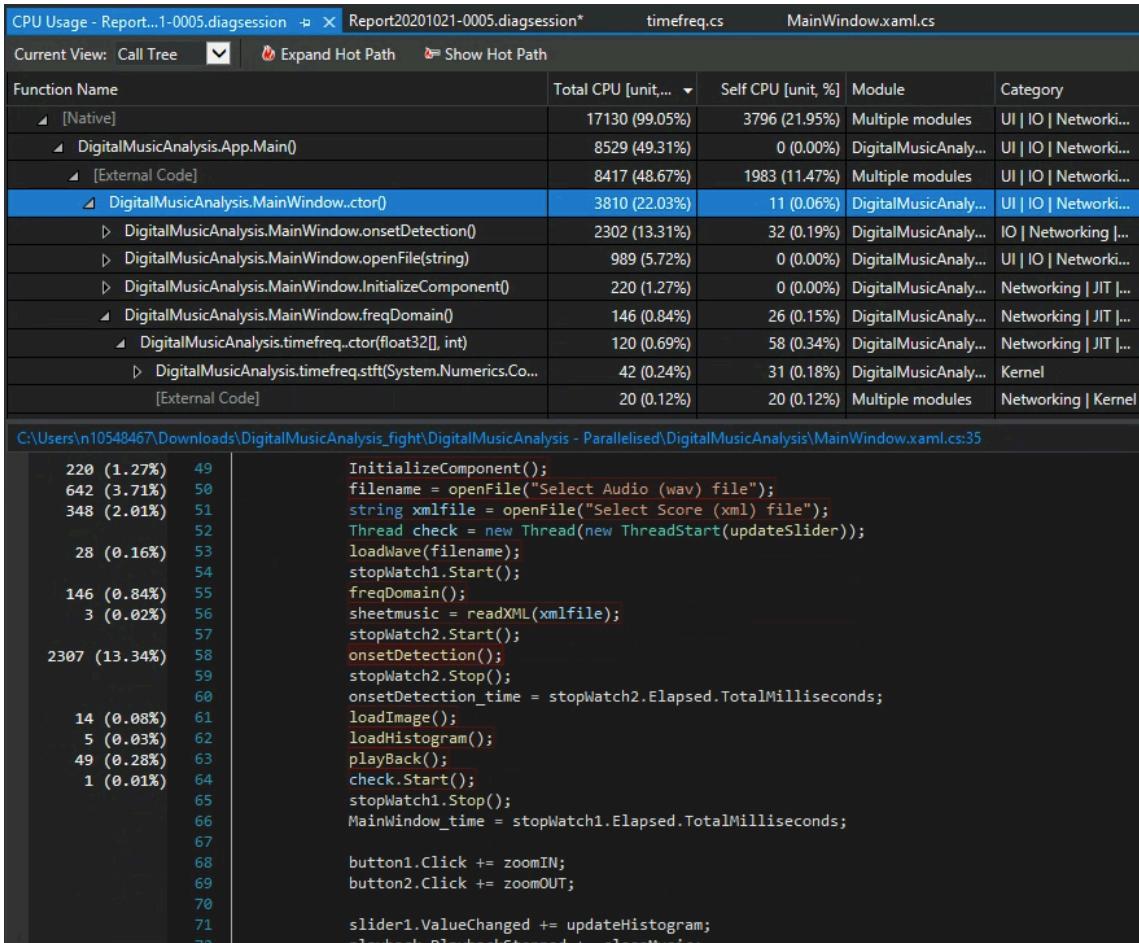


Fig.18 Profiling result of the parallelised version of the application

The time used in running each function is recorded in different level of parallelism exposure ranging from 1 to 8 processor counts. Moreover, a speedup ratio is calculated for each processor count of the MainWindow method as demonstrated in table 2.

MainWindow			stff		onsetDetection	
Processor count	time (ms)	speed-up	Processor count	time (ms)	Processor count	time (ms)
1	3272.7317	1.112943722	1	1484.1659	1	1540.8411
2	2602.9121	1.399342759	2	962.7147	2	1360.9276
3	2542.3967	1.432650617	3	855.7366	3	1388.0724
4	2425.2537	1.501849559	4	783.8631	4	1325.9001
5	2212.5913	1.646199278	5	698.0564	5	1208.3032
6	2190.7844	1.662585419	6	747.5258	6	1150.488
7	2363.2993	1.541220869	7	749.6073	7	1226.7276
8	2334.0133	1.560559316	8	814.3922	8	1208.1522

Table 1. Spent time on parallel version

After the speedup ratios are obtained, a speedup curve can be plotted as in the figure 19. It appears that the parallel version of the application can achieve slightly speedup

in a sublinear manner. This probably because despite having done a great job on the stff function, it is failed to enhance the performance of the onsetDetection function. Additionally, the stff seems still not scalable because after hit 5 processors, the time taken is stagnant and then increased.

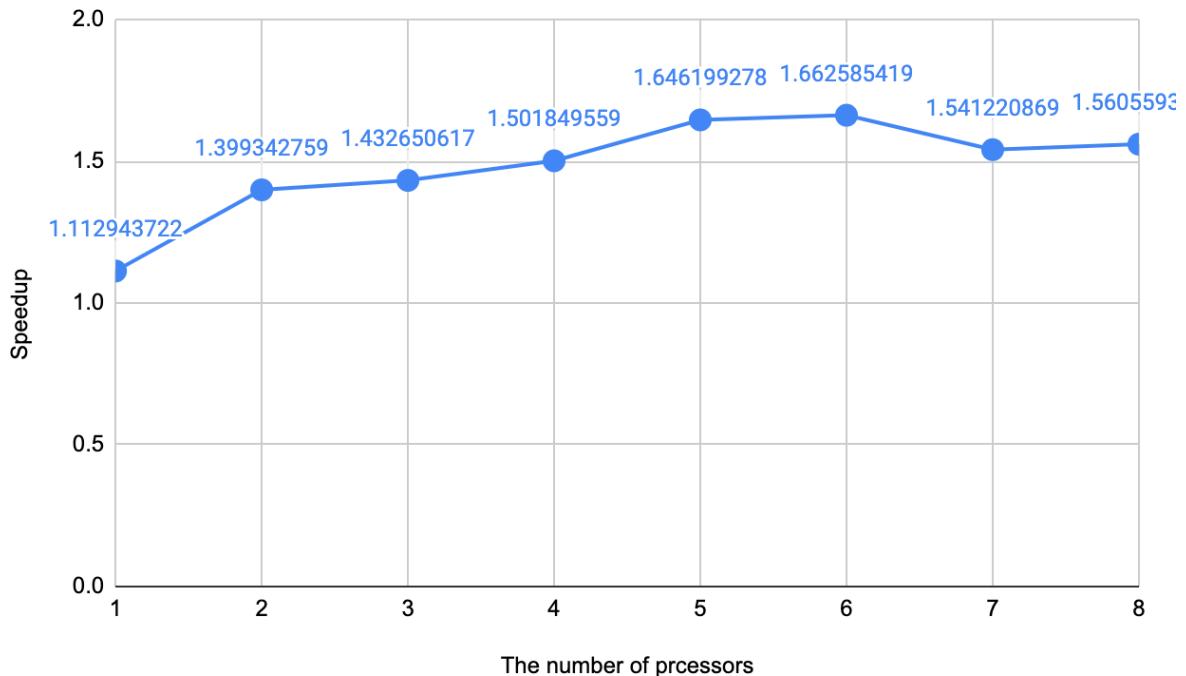


Fig.19 A speedup curve

4.3 Outputs

Since the outputs of the application are visualised by its nature, we can visually compare the outputs from the frequency and staff tabs of the program as in .

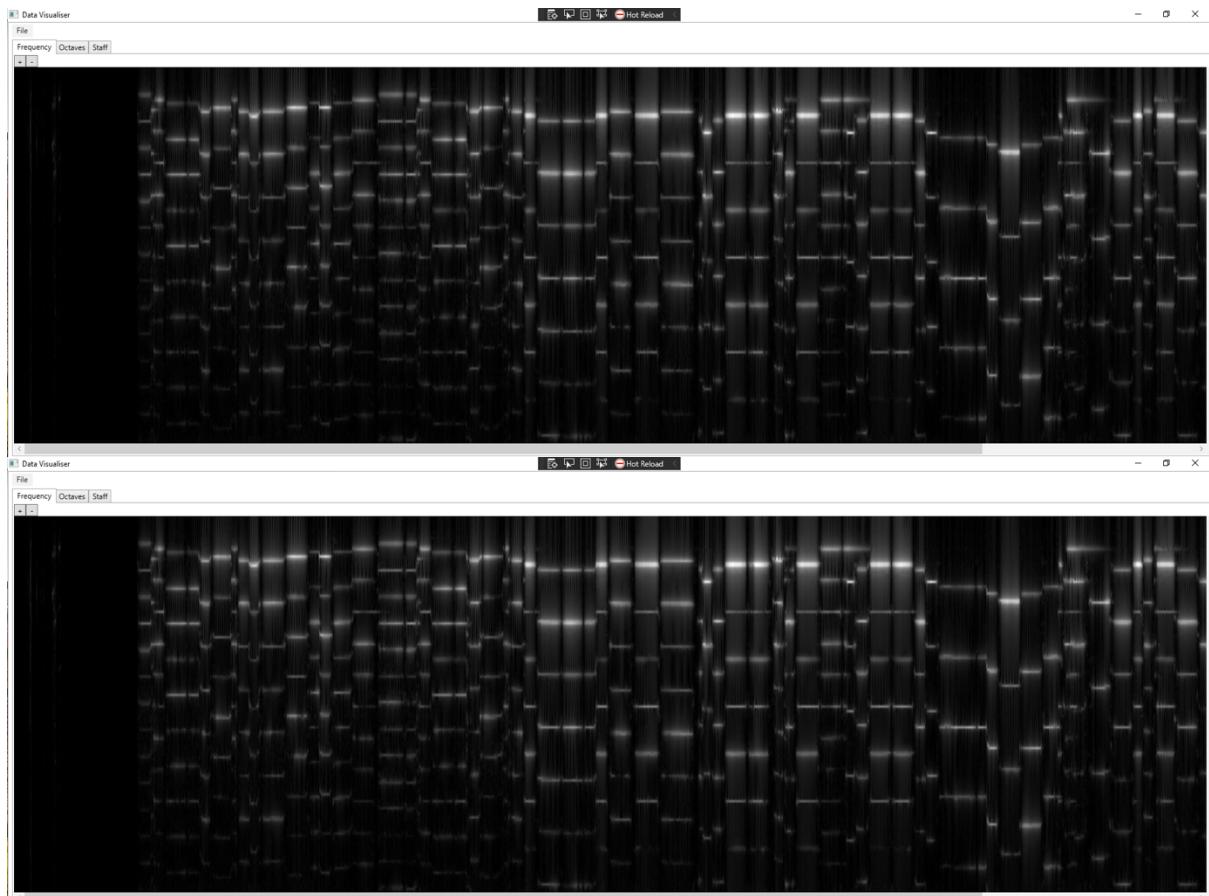
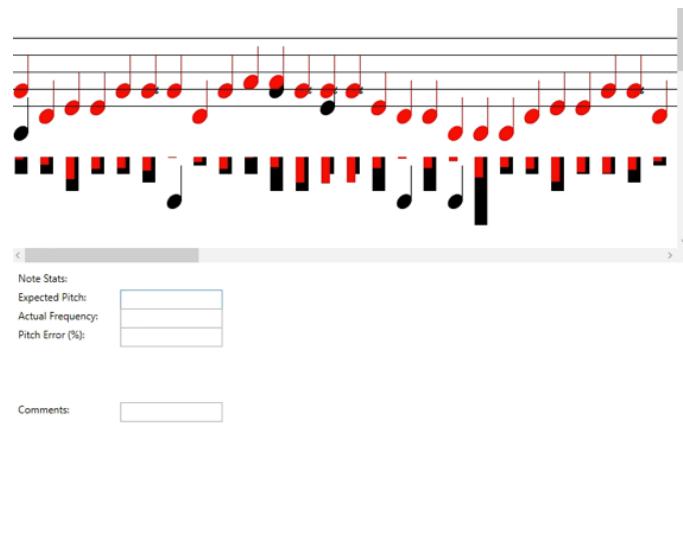


Fig. 20 Outputs of Frequency tab (upper: sequential version, lower: parallel version)



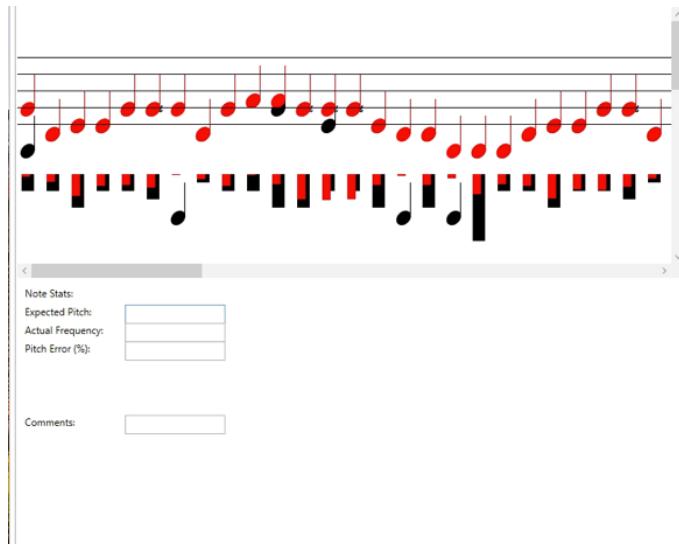


Fig. 21 Outputs of Staff tab (upper: sequential version, lower: parallel version)

5. Reflection

In this parallelisation assignment, I had an opportunity to enhance the performance of a substantial program application. Having look at the source-code at first, it was intimidating to see a number of lines and methods. Fortunately, there was a systematic approach to start solving a problem like this through using the Microsoft Visual Studio Profiler which I have never used before to grasp where the resource being spent substantially in specific sections of the program. This makes me think it will be very useful in my future career too. Subsequently, I was exposed to different ways to analyse variable dependencies as well as a number of ways to restructure a program to safely boost its performance from the lower-level abstraction instead of merely parallelise a loop. Above all, I was able to safely improve the performance of stff function by running in separated threads. However, I failed to increase speedup for the onsetDetection function as the dependencies could not be preserve so it results in slightly different in terms of execution time making the overall run time does not much improve. Finally, the future work is potentially changing from the current recursive FFT algorithm to be in an iterative form which can calculate each sample output.

6. Appendices

6.1 Appendix A

- Microsoft Task Parallel Library (TPL)
<https://docs.microsoft.com/en-us/dotnet/standard/parallel-programming/task-parallel-library-tpl>
- Parallel.For method
<https://docs.microsoft.com/en-us/dotnet/api/system.threading.tasks.parallel.for?view=netcore-3.1>
- The maximum number of concurrent tasks

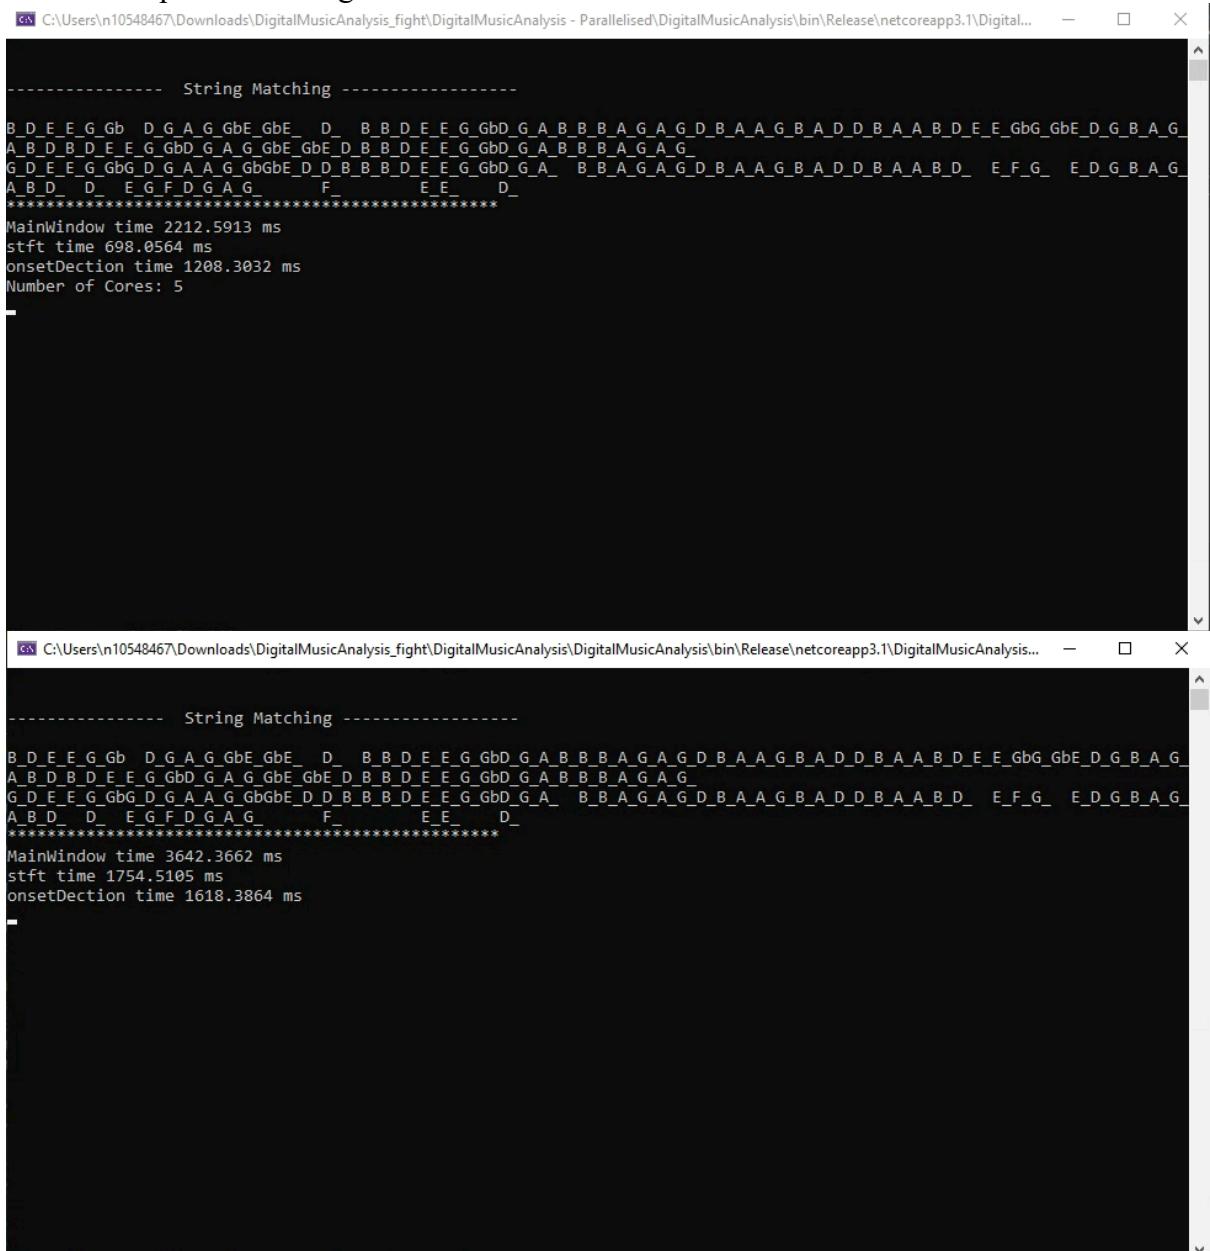
<https://docs.microsoft.com/en-us/dotnet/api/system.threading.tasks.paralleloptions.maxdegreeofparallelism?view=netcore-3.1>

- Stopwatch class

<https://docs.microsoft.com/en-us/dotnet/api/system.threading.tasks.parallel.for?view=netcore-3.1>

6.2 Appendix B

Console outputs describing the time taken in each function



The image shows two side-by-side screenshots of a terminal window. Both windows have the title bar "C:\Users\n10548467\Downloads\DigitalMusicAnalysis_fight\DigitalMusicAnalysis - Parallelised\DigitalMusicAnalysis\bin\Release\netcoreapp3.1\Digital...". The terminal is displaying the results of a "String Matching" operation. The output includes a large string of musical notes (B, D, E, F, G, A, C, E) repeated multiple times, followed by performance metrics: "MainWindow time 2212.5913 ms", "stft time 698.0564 ms", "onsetDetection time 12088.3032 ms", and "Number of Cores: 5". The second terminal window shows nearly identical output, indicating the same process was run again.

```
----- String Matching -----  
B_D_E_E_G_Gb_D_G_A_G_GbE_GbE_D_B_B_D_E_E_G_GbD_G_A_B_B_B_A_G_A_G_D_B_A_A_G_B_A_D_D_B_A_A_B_D_E_E_GbG_GbE_D_G_B_A_G  
A_B_D_B_D_E_E_G_GbD_G_A_G_GbE_GbE_D_B_B_D_E_E_G_GbD_G_A_B_B_B_A_G_A_G  
G_D_E_E_G_GbG_D_G_A_A_G_GbGbE_D_D_B_B_B_D_E_E_G_GbD_G_A_B_B_A_G_D_B_A_A_G_B_A_D_D_B_A_A_B_D_E_F_G_E_D_G_B_A_G  
A_B_D_D_E_G_F_D_G_A_G_F_E_E_D_*****  
MainWindow time 2212.5913 ms  
stft time 698.0564 ms  
onsetDetection time 12088.3032 ms  
Number of Cores: 5  
  
----- String Matching -----  
B_D_E_E_G_Gb_D_G_A_G_GbE_GbE_D_B_B_D_E_E_G_GbD_G_A_B_B_B_A_G_A_G_D_B_A_A_G_B_A_D_D_B_A_A_B_D_E_E_GbG_GbE_D_G_B_A_G  
A_B_D_B_D_E_E_G_GbD_G_A_G_GbE_GbE_D_B_B_D_E_E_G_GbD_G_A_B_B_B_A_G_A_G  
G_D_E_E_G_GbG_D_G_A_A_G_GbGbE_D_D_B_B_B_D_E_E_G_GbD_G_A_B_B_A_G_D_B_A_A_G_B_A_D_D_B_A_A_B_D_E_F_G_E_D_G_B_A_G  
A_B_D_D_E_G_F_D_G_A_G_F_E_E_D_*****  
MainWindow time 3642.3662 ms  
stft time 1754.5105 ms  
onsetDetection time 1618.3864 ms
```